



UsedCar-Price Prediction Project



Submitted by:
SARMISHTHA HALDAR

ACKNOWLEDGEMENT:

I would like to thank all my teachers, supervisors for the learning. Few journals referred in the case are as follows:

- i) Gareth, J., Daniela, W., Trevor, H., & Tibshirani, R. (2013). *An Introduction to Statistical*
- ii) Raschka, S., & Mirjalili, V. (2017). *Python machine learning*. Packt Publishing Ltd.

Business Problem Framing

This is a classic Business problem which helps to evaluate the price of the used car using the modelling below. The problem has occurred due to recent changes in the car market due to COVID-19 impact.

Conceptual Background

With COVID-19 impact in the market, we have seen lot of changes in the car market. Now some cars are in demand hence making them costly and some are not in demand and hence making them cheaper. With the change in market due to covid-19 impact the previous price evaluation models are not serving the purpose and hence we need to provide a car evaluation model which will help them to decide the car prices.

INTRODUCTION

Determining whether the listed price of a used car is a challenging task, due to the many factors that drive a used vehicle's price on the market. The focus of this project is developing machine learning models that can accurately predict the price of a used car based on its features, in order to make informed purchases. We implement and evaluate various learning methods on a dataset consisting of the sale prices of different makes and models across cities in India.

Problem Statement

To Build a model which can be used to predict prices of used cars

Analytical Problem Framing

We have used methods like r^2 score and RMSE for model evaluations

R^2 is a statistic that will give some information about the goodness of fit of a model. In regression, the R^2 **coefficient of determination** is a statistical measure of how well the regression predictions approximate the real data points. An R^2 of 1 indicates that the regression predictions perfectly fit the data

Root Mean Square Error (RMSE) is the standard deviation of the residuals (prediction errors). Residuals are a measure of how far from the regression line data points are; RMSE is a measure of how spread out these residuals are. In other words, it tells you how concentrated the data is around the line of best fit.

- **Data Sources and their formats**

We received the data in the form of .csv file and data was loaded using Pandas

```

1 [2]: Car=pd.read_csv('Cardetails.csv')

1 [3]: Car.head()

it[3]:

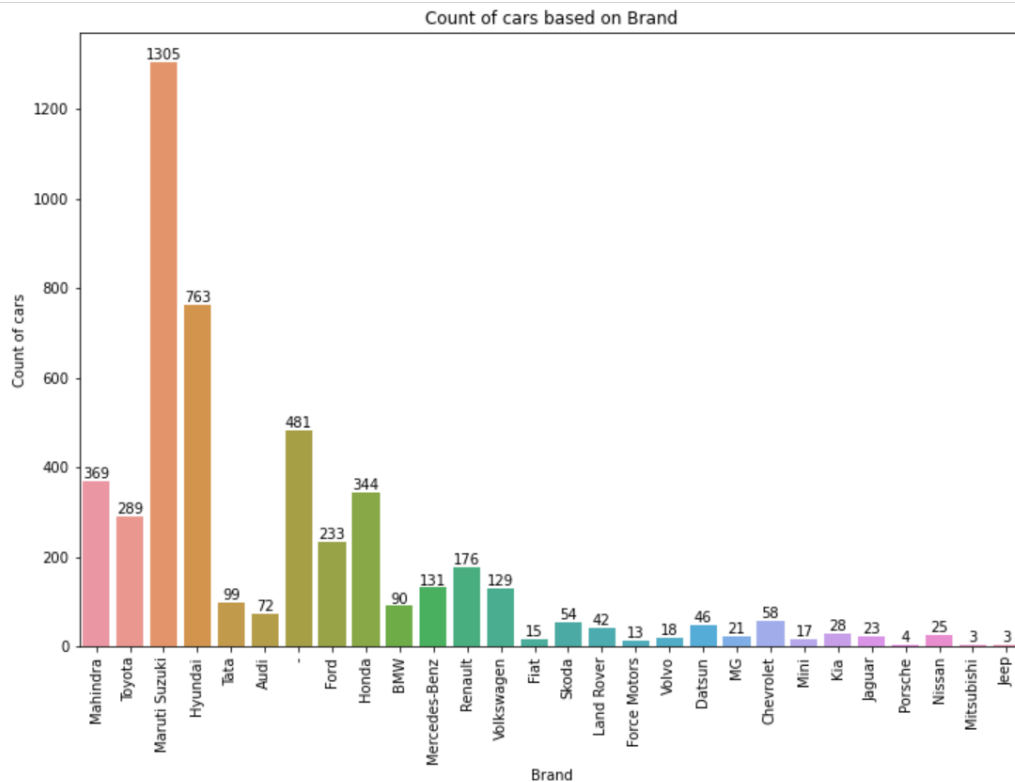
```

	Unnamed: 0	Brand	Model	Variant	Manufacturing_Year	Driven_kilometres	Fuel	Number_of_owners	Location	Price
0	0	Maruti Suzuki	Celerio	ZXI AMT	2017	11,439 km	Petrol	1st	Sainikpuri, Hyderabad, Telangana	₹ 4,90,000
1	1	Mahindra	Xylo	2009-2011 E8	2011	81,000 km	Diesel	1st	Rocktown Colony, Hyderabad, Telangana	₹ 3,96,000
2	2	Tata	Nexon	1.2 Revotron XM	2018	55,700 km	Petrol	1st	Himayat Nagar, Hyderabad, Telangana	₹ 7,75,000
3	3	Honda	CR-V	2007-2012 AT With Sun Roof	2010	71,174 km	Petrol	1st	Madhapur, Hyderabad, Telangana	₹ 7,50,000
4	4	Maruti Suzuki	Swift Dzire	VDI	2019	65,035 km	Diesel	1st	Ameerpet, Hyderabad, Telangana	₹ 7,90,000

- Exploratory Data Analysis
 - i)Brand

```
In [17]: Car['Brand'].value_counts()
```

```
Out[17]: Maruti Suzuki      2570  
         Hyundai          1090  
         Mahindra           530  
         Honda             478  
         Toyota            428  
         Ford              329  
         Renault           232  
         Mercedes-Benz      179  
         Volkswagen         175  
         Tata               162  
         BMW                129  
         Audi               93  
         Chevrolet          86  
         Skoda              74  
         Land Rover         72  
         Datsun             65  
         Kia                45  
         Jaguar             36  
         Volvo              31  
         Nissan             28  
         MG                 28  
         Fiat               20  
         Force Motors       18  
         Mini               18  
         Mitsubishi         6  
         Jeep               5  
         Porsche            4  
         Name: Brand, dtype: int64
```



Maximum cars in the dataset are by the manufacturer Maruti and looks like its quite popular .

- Location

```
In [19]: Car['Location'].value_counts()
Out[19]: -
Pitampura, Delhi, Delhi          706
Noida Extension, Noida, Uttar Pradesh  229
Madhapur, Hyderabad, Telangana    171
Hazratganj, Lucknow, Uttar Pradesh  146
...
Infocity, Gandhinagar, Gujarat      3
Subhash Park, Vadodara, Gujarat      2
Bhagwan Nagar Tekra, Ahmedabad, Gujarat  2
Billekahalli, Bengaluru, Karnataka    2
Banaswadi Rammurthi Nagar Green Park Layout, Bengaluru, Karnataka  2
Name: Location, Length: 260, dtype: int64
```

There are few missing values and we can look at it while data preprocessing

- Driven_kilometres

```

In [34]: X_train["Driven_kilometres"]

Out[34]: 3203      71,000 km
          1350     110,000 km
          6812      67,000 km
          446     175,835 km
          1743      61,000 km
          ...
          3772     15,000 km
          5191     90,000 km
          5226     39,000 km
          5390     39,000 km
          860      44,000 km
          Name: Driven_kilometres, Length: 4851, dtype: object

```

This clearly shows that data range is really high and high values might affect prediction thus it is important that scaling can be applied .

- **Year**

This simply displays the year which we have applied function to calculate the age of the car.

- **Data Preprocessing**

i)Removing the Unwanted columns

Remove the unwanted columns

```

In [4]: Car.drop('Unnamed: 0',inplace=True,axis=1)

In [5]: Car.head()

Out[5]:

```

	Brand	Model	Variant	Manufacturing_Year	Driven_kilometres	Fuel	Number_of_owners	Location	Price
0	Maruti Suzuki	Celerio	ZXI AMT	2017	11,439 km	Petrol	1st	Sainikpuri, Hyderabad, Telangana	₹ 4,90,000
1	Mahindra	Xylo	2009-2011 E8	2011	81,000 km	Diesel	1st	Rocktown Colony, Hyderabad, Telangana	₹ 3,96,000
2	Tata	Nexon	1.2 Revotron XM	2018	55,700 km	Petrol	1st	Himayat Nagar, Hyderabad, Telangana	₹ 7,75,000
3	Honda	CR-V	2007-2012 AT With Sun Roof	2010	71,174 km	Petrol	1st	Madhapur, Hyderabad, Telangana	₹ 7,50,000
4	Maruti Suzuki	Swift Dzire	VDI	2019	65,035 km	Diesel	1st	Ameerpet, Hyderabad, Telangana	₹ 7,90,000

```

In [6]: Car['Age'] = Car['Manufacturing_Year'].apply(lambda x: 2019 - x)

```

```
In [20]: Car['Location']=Car['Location'].replace('-', 'Delhi')
```

```
In [21]: Car.Location.mode()
```

```
Out[21]: 0    Delhi  
dtype: object
```

Location should not be a determinant for the price of a car and I'll safely remove it.

ii) Checking for null/missing values and imputing them with mean,mode, median as required

i)Dealing with missing values

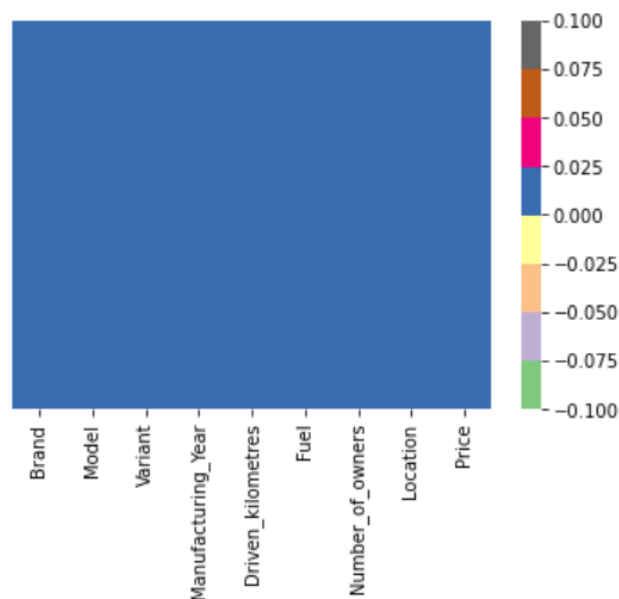
```
In [35]: Car.isnull().sum()
```

```
Out[35]: Brand          0  
Model          0  
Variant        0  
Manufacturing_Year  0  
Driven_kilometres  0  
Fuel           0  
Number_of_owners  0  
Location       0  
Price          0  
dtype: int64
```

```
In [36]: #There are no missing values in the dataframe
```

```
In [37]: #heatmap to verify nulll values using graph  
sns.heatmap(Car.isnull(),yticklabels=False,cbar=True,cmap='Accent')
```

```
Out[37]: <matplotlib.axes._subplots.AxesSubplot at 0x265bf665280>
```



iii) Converting the columns to required data types and meaningful data

Year

```
In [24]: curr_time = datetime.datetime.now()
```

```
In [31]: X_train['Manufacturing_Year']=X_train['Manufacturing_Year'].astype(int)
X_test['Manufacturing_Year'] = X_test['Manufacturing_Year'].astype(int)
```

```
In [32]: X_train['Manufacturing_Year'] = X_train['Manufacturing_Year'].apply(lambda x : curr_time.year - x)
X_test['Manufacturing_Year'] = X_test['Manufacturing_Year'].apply(lambda x : curr_time.year - x)
```

```
In [33]: X_train['Manufacturing_Year']
```

```
Out[33]: 3203    9
1350    7
6812   10
446     7
1743    9
      ..
3772    3
5191   11
5226   20
5390   10
860     5
Name: Manufacturing_Year, Length: 4851, dtype: int64
```

Model,Variant,Fuel,Number_of_owners .All these columns are categorical columns which should be converted to dummy variables before being used

```
In [38]: #define numeric variable and categorical variable to work separatly on them
num_col=['Manufacturing_Year']
cat_cols=['Brand','Model','Variant','Fuel','Number_of_owners','Location','Driven_kilometres']
```

Now that we have worked with the training data, let's create dummy columns for categorical columns before we begin training.

```
In [39]: X_train = pd.get_dummies(X_train,
                                columns = ["Brand", "Model", "Variant","Fuel", "Number_of_owners","Driven_kilometres"],
                                drop_first = True)
```

```
In [40]: X_test = pd.get_dummies(X_test,
                                columns = ["Brand", "Model", "Variant","Fuel", "Number_of_owners","Driven_kilometres"],
                                drop_first = True)
```

```
In [41]: #It might be possible
#that the dummy column creation would be different in test and train data, thus, I'd fill in all missing columns with zeros.
```

```
In [42]: missing_cols = set(X_train.columns) - set(X_test.columns)
for col in missing_cols:
    X_test[col] = 0
X_test = X_test[X_train.columns]
```

Iv) Scaling the data

```
In [43]: standardScaler = StandardScaler()
```

```
In [44]: standardScaler.fit(X_train)
```

```
Out[44]: StandardScaler()
```

```
In [45]: X_train = standardScaler.transform(X_train)
X_test = standardScaler.transform(X_test)
```

```
- . . . . .
```

- Hardware and Software Requirements , Tools Used
No Specific requirements except Jupyter Notebook.

Model/s Development and Evaluation

- Identification of possible problem
This is a Regression Problem and we have used Decision tree, Random Forest, KNN , LASSO, XGboost, Gradientboosting regressor to build the model.
- Run and Evaluate selected models: We used Linear regression, Decision tree regressor , Random forest Regressor and found that Decision tree and random forest are performing well when we found the r^2 score. Then we further did a gridsearch tuning with host of other models and then created an average of the best performing models.

- **KNNRidge regression:**

KNN has been used in **statistical estimation and pattern recognition** already in the beginning of 1970's as a non-parametric technique. A simple implementation of KNN regression is to calculate the average of the numerical target of the K nearest neighbors

```
In [56]: KRR = KernelRidge()

KRR_grid = {"alpha" : [25,10,4,2,1.0,0.8,0.5,0.3,0.2,0.1,0.05,0.02,0.01],
            "kernel" : ["polynomial"],
            "degree" : [1,2,3,4,5],
            "coef0" : [1,1.5,2,2.5,3,3.5,4,4.5,5]
            }

KRRModel = GridSearchCV(estimator = KRR, param_grid = KRR_grid, cv=kf, scoring="neg_mean_squared_error", n_jobs= 4, verbose = 1)
KRRModel.fit(X_train,y_train)
KRR_best = KRRModel.best_estimator_
KRRModel.best_params_

Fitting 5 folds for each of 585 candidates, totalling 2925 fits
Out[56]: {'alpha': 0.01, 'coef0': 3, 'degree': 2, 'kernel': 'polynomial'}
```

Gradient boosting regressor: GradientBoostingRegressor GB builds an additive model in a forward stage-wise fashion; it allows for the optimization of arbitrary differentiable loss functions. In each stage a regression tree is fit on the negative gradient of the given loss function.

```
In [66]: GBR = GradientBoostingRegressor()

GBR_grid = {"n_estimators" : [2000,3000],
            "learning_rate" : [0.01,0.1],
            "max_depth" : [3,5],
            "max_features" : ['sqrt'],
            "min_samples_leaf" : [10,15],
            "min_samples_split" : [2,5],
            "loss" : ['huber']
            }

GBRModel = GridSearchCV(estimator = GBR, param_grid = GBR_grid, cv=kf, scoring="neg_mean_squared_error", n_jobs= 4, verbose = 1)
GBRModel.fit(train,y_train)
GBR_best = GBRModel.best_estimator_
GBRModel.best_params_

Fitting 5 folds for each of 32 candidates, totalling 160 fits
Out[66]: {'learning_rate': 0.01,
          'loss': 'huber',
          'max_depth': 3,
          'max_features': 'sqrt',
          'min_samples_leaf': 15,
          'min_samples_split': 5,
          'n_estimators': 3000}
```

XGBRegressor: It is an efficient implementation of gradient boosting that can be used for regression predictive modeling

```
In [67]: XGB = XGBRegressor()

XGB_grid = {'nthread':[4],
            'objective':['reg:linear'],
            'learning_rate': [.03, 0.05, .07],
            'max_depth': [5, 6, 7],
            'min_child_weight': [4],
            'silent': [1],
            'subsample': [0.7],
            'colsample_bytree': [0.7],
            'n_estimators': [500]}

XGBModel = GridSearchCV(estimator = XGB, param_grid = XGB_grid, cv=kf, scoring="neg_mean_squared_error", n_jobs= 4, verbose = 1)
XGBModel.fit(train,y_train)
XGB_best = XGBModel.best_estimator_
XGBModel.best_params_
```

```
Out[67]: {'colsample_bytree': 0.7,
          'learning_rate': 0.07,
          'max_depth': 5,
          'min_child_weight': 4,
          'n_estimators': 500,
          'nthread': 4,
          'objective': 'reg:linear',
          'silent': 1,
          'subsample': 0.7}
```

Decision Tree Regressor: Decision Tree - Regression. Decision tree builds regression or classification models in the form of a tree structure. **It breaks down a dataset into smaller and smaller subsets while at the same time an associated decision tree is incrementally developed.** The final result is a tree with decision nodes and leaf nodes.

```
In [59]: DTR = DecisionTreeRegressor(random_state=0)

DTR_grid = {
    "criterion": ['mse'],
    "splitter": ['best'],
    "max_depth": [2,3,5,10],
    "max_features": ['sqrt'],
    "min_samples_leaf": [5,10,15],
    "min_samples_split": [1,2,5]
}

DTRModel = GridSearchCV(estimator = DTR, param_grid = DTR_grid, cv=kf, scoring="neg_mean_squared_error", n_jobs= 4, verbose = 1)
DTRModel.fit(X_train,y_train)
DTR_best = DTRModel.best_estimator_
DTRModel.best_params_

Fitting 5 folds for each of 36 candidates, totalling 180 fits

Out[59]: {'criterion': 'mse',
          'max_depth': 10,
          'max_features': 'sqrt',
          'min_samples_leaf': 5,
          'min_samples_split': 2,
          'splitter': 'best'}
```

LASSO: Lasso regression is a **regularization technique**. It is used over regression methods for a more accurate prediction. This model uses shrinkage. Shrinkage is where data values are shrunk towards a central point as the mean. The lasso procedure encourages simple, sparse models (i.e. models with fewer parameters).

```
In [60]: LARS = Lasso()

LARS_grid = {
    "alpha": [1,0.8,0.3,0.2,0.1,0.05,0.005,0.02,0.01],
    "max_iter": [500,700,1000]
}

LARSModel = GridSearchCV(estimator = LARS, param_grid = LARS_grid, cv=kf, scoring="neg_mean_squared_error", n_jobs= 4, verbose = 1)
LARSModel.fit(X_train,y_train)
LARS_best = LARSModel.best_estimator_
LARSModel.best_params_

Fitting 5 folds for each of 27 candidates, totalling 135 fits

Out[60]: {'alpha': 1, 'max_iter': 1000}
```

```

In [61]: n_folds = 5

def rmsle_cv(model):
    kf = KFold(n_folds, shuffle=True, random_state=42).get_n_splits(X_train)
    rmse= np.sqrt(-cross_val_score(model, X_train, y_train, scoring="neg_mean_squared_error", cv = kf))
    return(rmse)

In [62]: #Build the base models based on GridSearch tuning

KRR = KernelRidge(alpha=0.8, coef0=5, degree=2, gamma=None, kernel='polynomial', kernel_params=None)

lasso = make_pipeline(RobustScaler(), Lasso(alpha =0.0005, max_iter = 500, random_state=1))

GBoost = GradientBoostingRegressor(n_estimators=3000, learning_rate=0.01, max_depth=3, max_features='sqrt', min_samples_leaf=10, min_samples_split=5, loss='huber')

model_xgb = xgb.XGBRegressor(colsample_bytree=0.7, learning_rate=0.03, max_depth=6, min_child_weight=4, n_estimators=500, subsample=0.7, silent=1, random_state =7)

model_DTR = DecisionTreeRegressor(criterion = 'mse', max_depth = 10, max_features = 'sqrt', min_samples_leaf = 5, min_samples_split = 2, splitter = 'best' )

LRModel = LinearRegression()

In [64]: score = rmsle_cv(KRR)
print("Kernel Ridge score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))

score = rmsle_cv(lasso)
print("Lasso score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))

score = rmsle_cv(GBoost)
print("Gradient Boosting score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))

score = rmsle_cv(model_xgb)
print("Xgboost score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))

score = rmsle_cv(model_DTR)
print("DT Regression score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))

score = rmsle_cv(LRModel)
print("Linear Regression score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))

Kernel Ridge score: 36216.4901 (12173.1580)

Lasso score: 29603.2268 (9415.6308)

Gradient Boosting score: 359391.7890 (30830.6505)

Xgboost score: 131510.8550 (13899.1997)

DT Regression score: 644898.8933 (45017.9086)

Linear Regression score: 805740048700595200.0000 (330606285470877888.0000)

```

Averaging the best models to create a model

```
In [65]: class AveragingModels(BaseEstimator, RegressorMixin, TransformerMixin):
def __init__(self, models):
    self.models = models

    # we define clones of the original models to fit the data in
def fit(self, X, y):
    self.models_ = [clone(x) for x in self.models]

    # Train cloned base models
    for model in self.models_:
        model.fit(X, y)

    return self

    #Now we do the predictions for cloned models and average them
def predict(self, X):
    predictions = np.column_stack([
        model.predict(X) for model in self.models_
    ])
    return np.mean(predictions, axis=1)
```

```
In [66]: #Averaging the best models to optimize the prediction.

#averaged_models = AveragingModels(models = (KRR, Lasso, ENet, GBoost, model_xgb, model_lgb, model_DTR, LRModel))
averaged_models = AveragingModels(models = (KRR,GBoost,model_DTR, lasso, model_xgb))

score = rmsle_cv(averaged_models)
print(" Averaged base models score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))
```

Averaged base models score: 210733.5443 (16892.4285)

```
]: averaged_models.fit(X_train, y_train)
averaged_models_pred = np.expml(averaged_models.predict(X_test))
```

- Key Metrics for success in solving problem under consideration

```
In [65]: class AveragingModels(BaseEstimator, RegressorMixin, TransformerMixin):
def __init__(self, models):
    self.models = models

    # we define clones of the original models to fit the data in
def fit(self, X, y):
    self.models_ = [clone(x) for x in self.models]

    # Train cloned base models
    for model in self.models_:
        model.fit(X, y)

    return self

    #Now we do the predictions for cloned models and average them
def predict(self, X):
    predictions = np.column_stack([
        model.predict(X) for model in self.models_
    ])
    return np.mean(predictions, axis=1)
```

```
In [66]: #Averaging the best models to optimize the prediction.

#averaged_models = AveragingModels(models = (KRR, Lasso, ENet, GBoost, model_xgb, model_lgb, model_DTR, LRModel))
averaged_models = AveragingModels(models = (KRR,GBoost,model_DTR, lasso, model_xgb))

score = rmsle_cv(averaged_models)
print(" Averaged base models score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))
```

Averaged base models score: 210733.5443 (16892.4285)

```
]: averaged_models.fit(X_train, y_train)
averaged_models_pred = np.expml(averaged_models.predict(X_test))
```

CONCLUSION

- **Conclusions**

According to the performance metrics, Lasso, Decisiontree regressor, KNN, Gradient boosting have good scores and hence we build a model by averaging the best ones as above.

- **Limitations & Scope for Future**

This study used different models in order to predict used car prices. However, there was a relatively small dataset for making a strong inference because number of observations was only 7000. Gathering more data can yield more robust predictions. Secondly, there could be more features that can be good predictors. For example, here are some variables that might improve the model: number of doors, gas/mile (per gallon), color, mechanical and cosmeticreconditioning time, used-to-new ratio, appraisal-to-trade ratio.

Another point that that has room to improvement is that data cleaning process can be done more rigorously with the help of more technical information